


Neural networks

Drawing computation graphs



what **expression** does this compute?

equivalently, what **program** does this correspond to?

$$|| (x_1\theta_1 + x_2\theta_2) - y ||^2$$

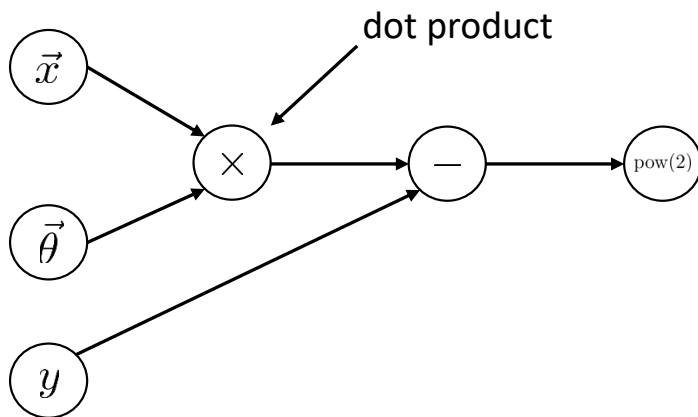
this is a **MSE loss** with a **linear regression** model

neural networks are **computation graphs**

if we design **generic tools** for computation graphs, we can train **many kinds** of neural networks

Drawing computation graphs

a simpler way to draw the same thing:



I'll drop the $\vec{}$ decorator from now on...

what **expression** does this compute?

equivalently, what **program** does this correspond to?

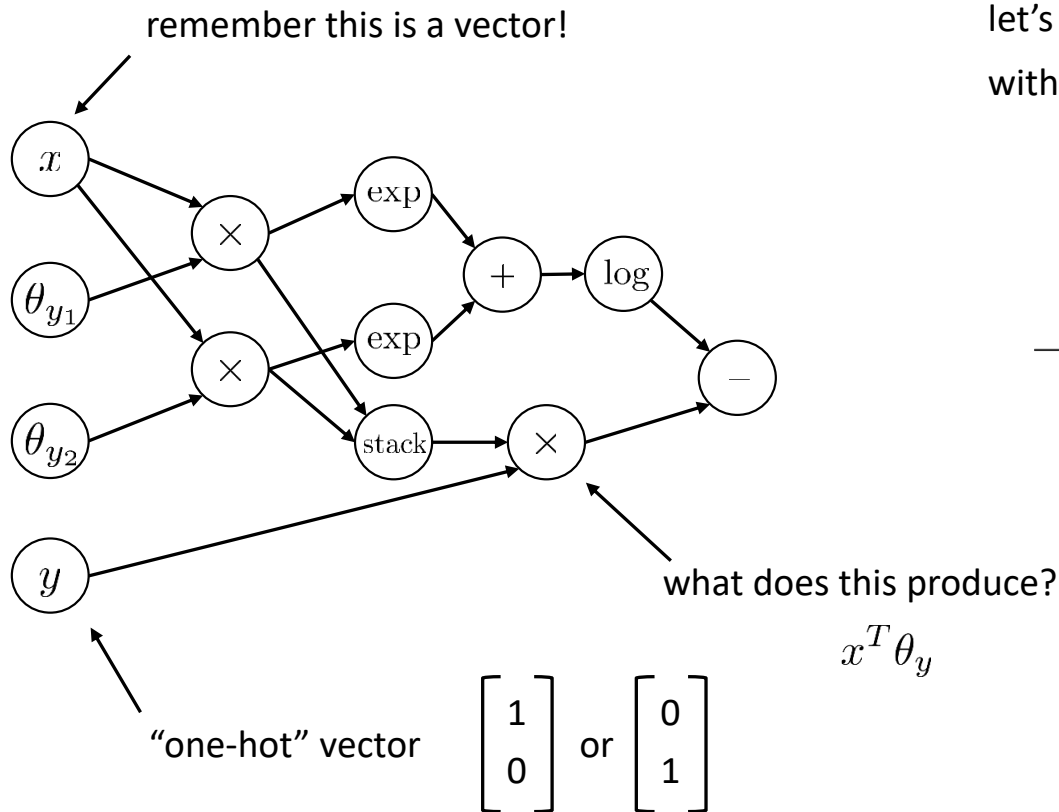
$$|| (x_1\theta_1 + x_2\theta_2) - y ||^2$$

this is a **MSE loss** with a **linear regression** model

neural networks are **computation graphs**

if we design **generic tools** for computation graphs, we
can train **many kinds** of neural networks

Logistic regression



let's draw the computation graph for **logistic regression**
with the negative log-likelihood loss

cục này là cục softmax

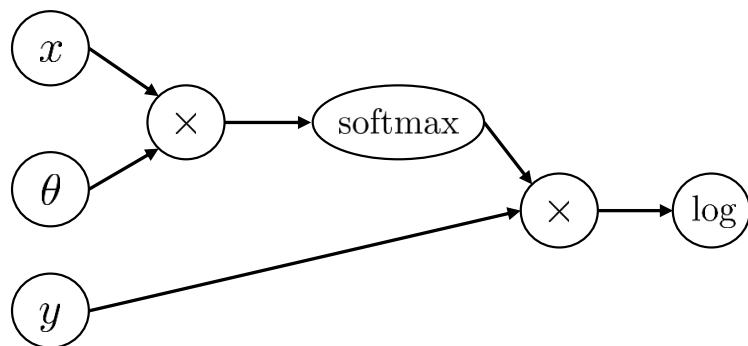
$$p_{\theta}(y|x) = \frac{\exp(x^T \theta_y)}{\sum_{y'} \exp(x^T \theta_{y'})}$$

$$-\log p_{\theta}(y|x) = -x^T \theta_y + \log \sum_{y'} \exp(x^T \theta_{y'})$$

Logistic regression

a simpler way to draw the same thing:

$$f_{\theta}(x) = \begin{bmatrix} x^T \theta_{y_1} \\ x^T \theta_{y_2} \\ \vdots \\ x^T \theta_{y_m} \end{bmatrix} \quad f_{\theta}(x) = \underset{\substack{\uparrow \\ \text{matrix}}}{\theta} x$$



$$p_{\theta}(y|x) = \frac{\exp(x^T \theta_y)}{\sum_{y'} \exp(x^T \theta_{y'})}$$

$$-\log p_{\theta}(y|x) = -x^T \theta_y + \log \sum_{y'} \exp(x^T \theta_{y'})$$

$$\begin{bmatrix} \theta_{y_1} \\ \theta_{y_2} \\ \theta_{y_3} \end{bmatrix} \times \begin{bmatrix} x \end{bmatrix} = \begin{bmatrix} x^T \theta_{y_1} \\ x^T \theta_{y_2} \\ \vdots \\ x^T \theta_{y_m} \end{bmatrix}$$

$$p_{\theta}(y = i|x) = \text{softmax}(f_{\theta}(x))[i] = \frac{\exp(f_{\theta,i}(x))}{\sum_{j=1}^m \exp(f_{\theta,j}(x))}$$

Drawing it even *more* concisely

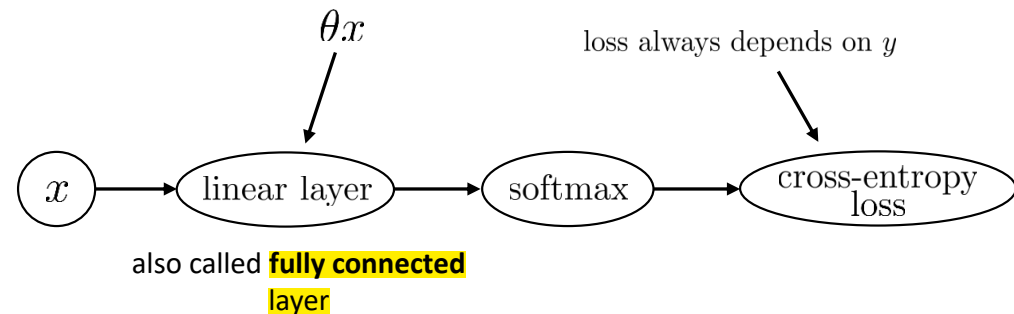
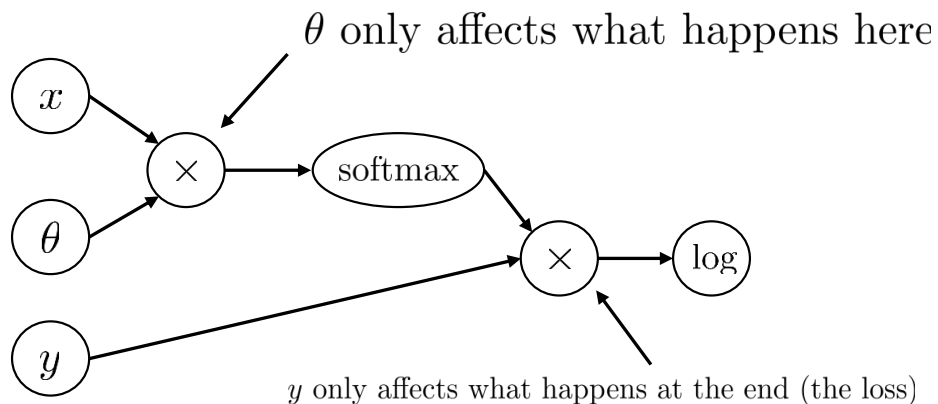
Notice that we have **two types** of variables:

data (e.g., x, y), which serves as input or target output

parameters (e.g., θ)


the parameters *usually* affect one specific operation

(though there is often *parameter sharing*, e.g., conv nets – more on this later)

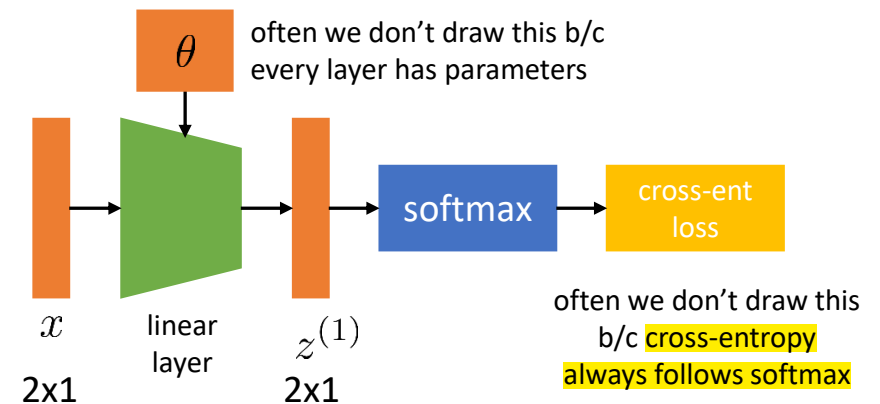


Neural network diagrams


(simplified) computation graph diagram



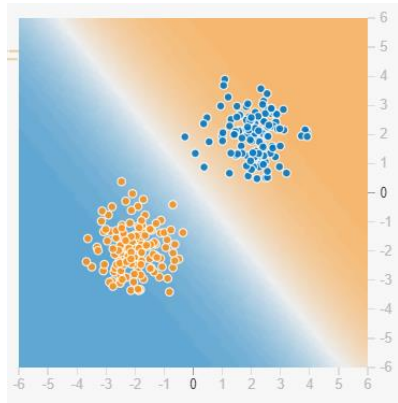
neural network diagram



simplified drawing:

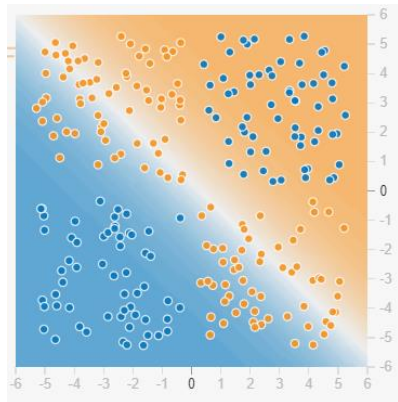


Logistic regression with features



$$\text{softmax}(x^T \theta)$$

pop quiz: what is the dimensionality of θ ?



$$\text{softmax}(\phi(x)^T \theta)$$

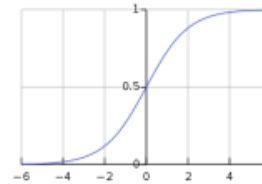
$$\phi(x) = \begin{pmatrix} x_1 \\ x_2 \\ x_1^2 \\ x_2^2 \\ x_1 x_2 \end{pmatrix}$$

Learning the features

Problem: how do we represent the learned features?

Idea: what if each feature is a (binary) logistic regression output?

$$\phi_1(x) = \text{softmax}(x^T w_1^{(1)}) = \frac{1}{1 + \exp(-x^T w_1^{(1)})}$$



$w_1^{(1)}$ ← which layer
← which feature
= rows of weight **matrix**

$$W^{(1)} = \begin{bmatrix} w_1^{(1)} \\ w_2^{(1)} \\ w_3^{(1)} \end{bmatrix}$$

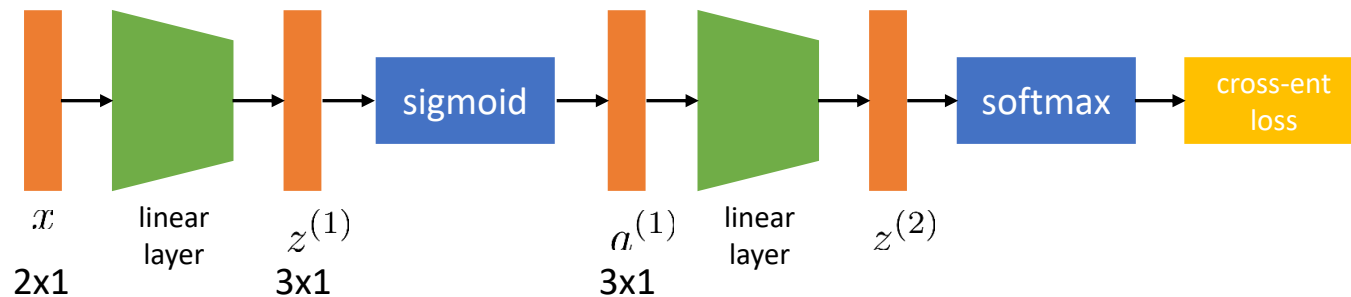
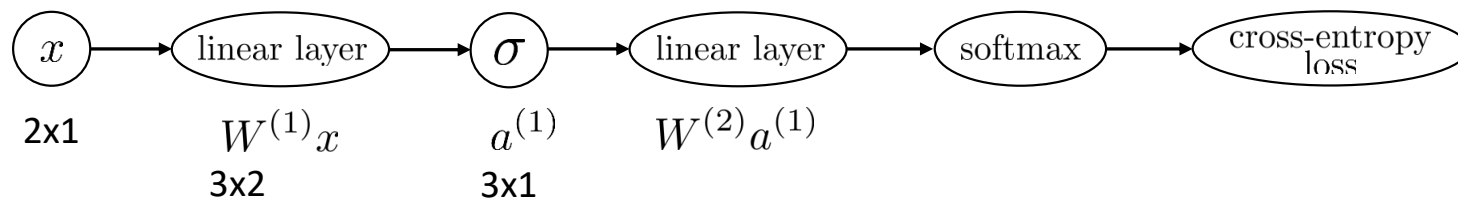
$$\phi(x) = \begin{pmatrix} \text{softmax}(x^T w_1^{(1)}) \\ \text{softmax}(x^T w_2^{(1)}) \\ \text{softmax}(x^T w_3^{(1)}) \end{pmatrix} = \sigma(W^{(1)}x)$$

aside: I'll switch to use w or W instead of θ here
 θ – all parameters of the model
 $w_1^{(1)}$ – weights (a.k.a. parameters) of feature 1 at layer 1


per-element sigmoid
not the same as softmax
each feature is independent

Let's draw this!


$$\phi(x) = \begin{pmatrix} \text{softmax}(x^T w_1^{(1)}) \\ \text{softmax}(x^T w_2^{(1)}) \\ \text{softmax}(x^T w_3^{(1)}) \end{pmatrix} = \sigma(W^{(1)}x) \quad p(y|x) = \text{softmax}(\phi(x)^T \theta)$$




Simpler drawing





simpler way to draw the same thing:



even simpler:




Doing it multiple times



Activation functions

$$\phi_1(x) = \text{softmax}(x^T w_1^{(1)}) = \frac{1}{1 + \exp(-x^T w_1^{(1)})}$$



we **don't have to** use a **sigmoid**!

a wide range of non-linear functions will work

these are called **activation functions**

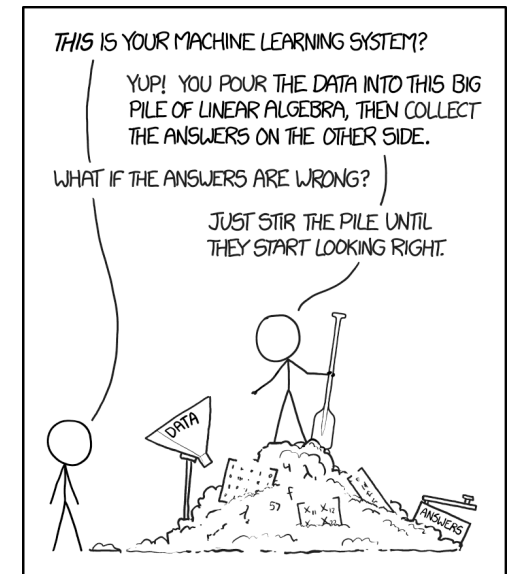
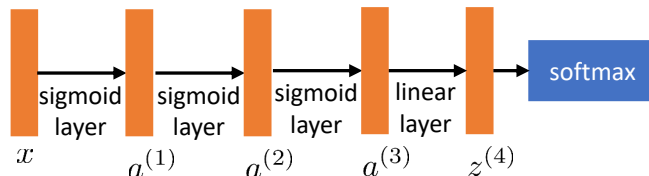
we'll discuss specific choices later

why **non-linear**?

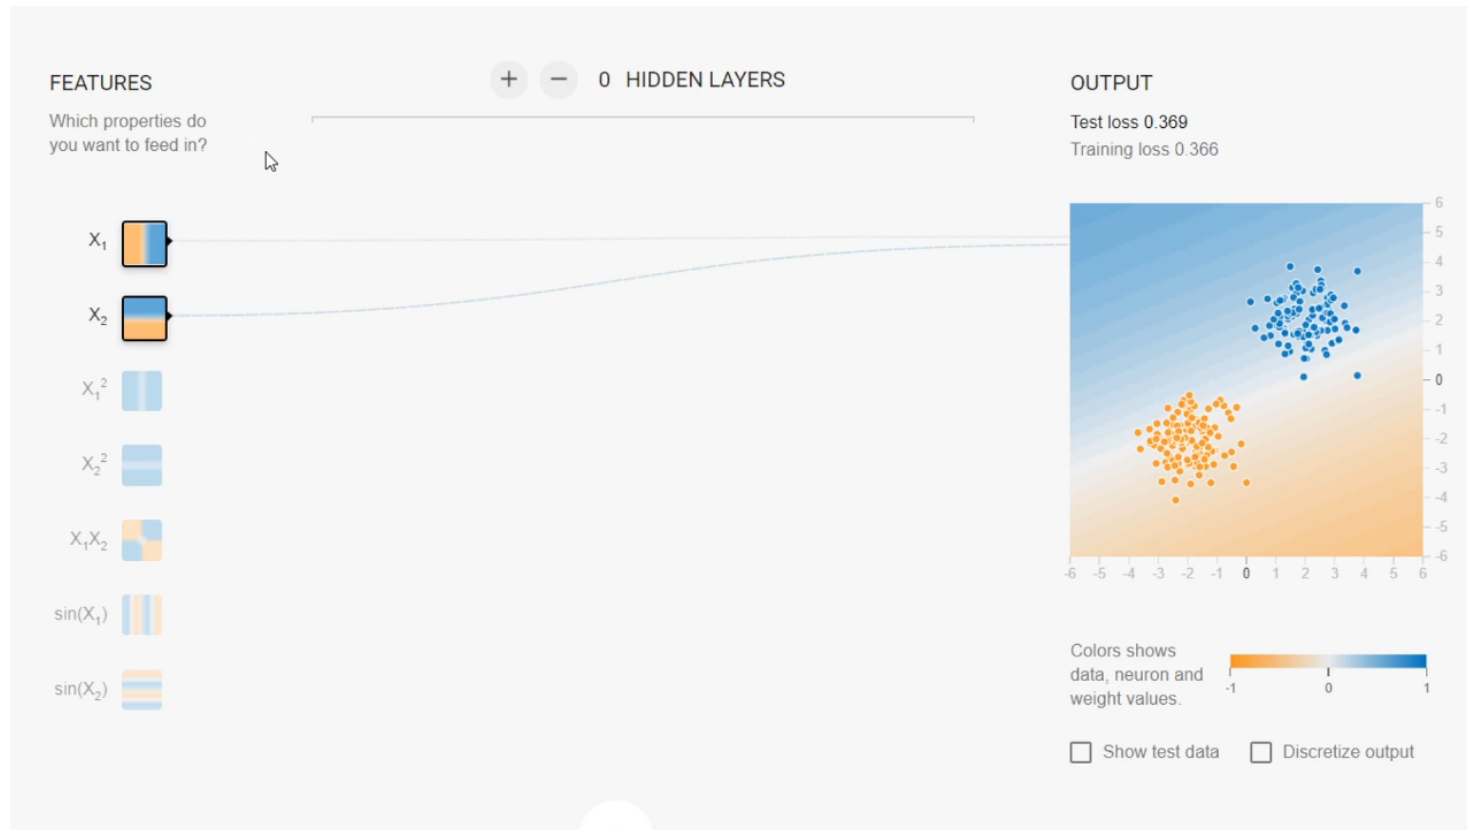
$$a^{(2)} = \sigma(W^{(2)}\sigma(W^{(1)}x)) \quad \text{if } \sigma(z) = z, \text{ then...}$$
$$a^{(2)} = W^{(2)}W^{(1)}x = Mx$$

multiple linear layers = one linear layer

enough layers = we can represent anything (so long as they're nonlinear)



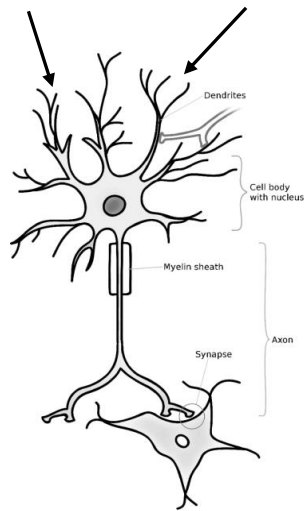
Demo time!



Source: <https://playground.tensorflow.org/>

Aside: what's so neural about it?


dendrites receive signals from other neurons



neuron "decides" whether to fire based on incoming signals

axon transmits signal to downstream neurons

artificial "neuron" sums up signals from upstream neurons (also referred to as "units")



neuron "decides" how much to fire based on incoming signals

activations transmitted to downstream units

$$z = \sum_i a_i$$

upstream activations

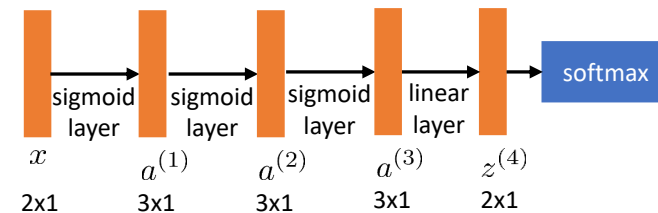
$$a = \sigma(z)$$

activation function

Training neural networks

What do we need?

1. Define your **model class**



2. Define your **loss function**

negative log-likelihood, just like before

3. Pick your **optimizer**

stochastic gradient descent
what do we need?

4. Run it on a big GPU

$$\nabla_{\theta} \mathcal{L}(\theta) = \begin{pmatrix} \frac{d\mathcal{L}(\theta)}{d\theta_1} \\ \frac{d\mathcal{L}(\theta)}{d\theta_2} \\ \vdots \\ \frac{d\mathcal{L}(\theta)}{d\theta_n} \end{pmatrix}$$

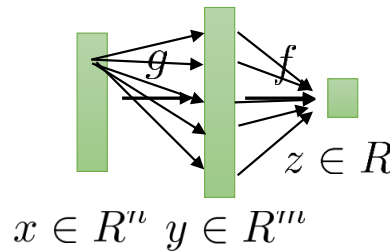
Aside: chain rule

Chain rule: $x \xrightarrow{g} y \xrightarrow{f} z$

$$\frac{d}{dx} f(g(x)) = \frac{dz}{dx} = \frac{dy}{dx} \frac{dz}{dy}$$


Jacobian of g

Jacobian of f




Row or column?

In this lecture:



In some textbooks:



High-dimensional chain rule

$$\frac{d}{dx_i} f(g(x)) = \sum_{j=1}^m \frac{dy_j}{dx_i} \frac{dz}{dy_j} = \frac{dy}{dx_i} \frac{dz}{dy}$$

\uparrow \uparrow \uparrow
 sum over all dimensions of y row $1 \times m$ col $m \times 1$

$$\frac{d}{dx} f(g(x)) = \frac{dy}{dx} \frac{dz}{dy}$$

\uparrow \uparrow
 mat $n \times m$ col $m \times 1$

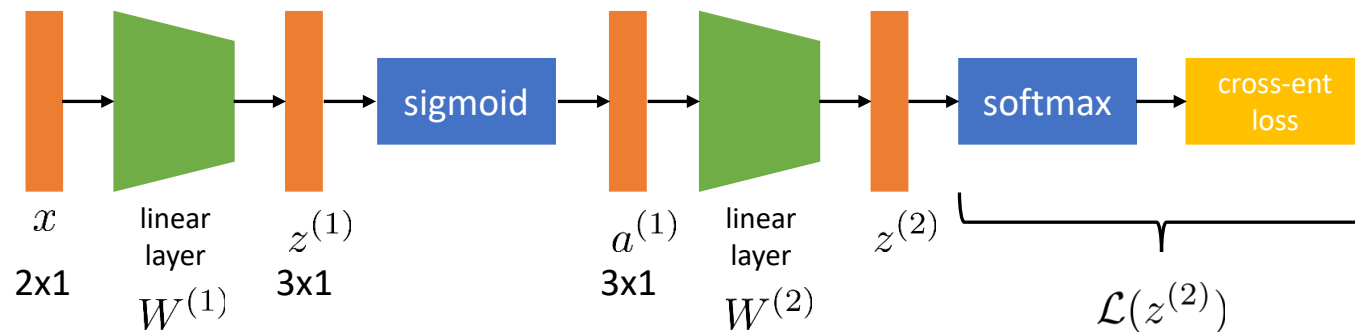
$$\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$$

Just two different conventions!

Chain rule for neural networks

A neural network is just a composition of functions

So we can use chain rule to compute gradients!



$$\frac{d\mathcal{L}}{dW^{(1)}} = \frac{dz^{(1)}}{dW^{(1)}} \frac{da^{(1)}}{dz^{(1)}} \frac{dz^{(2)}}{da^{(1)}} \frac{d\mathcal{L}}{dz^{(2)}}$$

$$\frac{d\mathcal{L}}{dW^{(2)}} = \frac{dz^{(2)}}{dW^{(2)}} \frac{d\mathcal{L}}{dz^{(2)}}$$

Does it work?

$$\frac{d\mathcal{L}}{dW^{(1)}} = \frac{dz^{(1)}}{dW^{(1)}} \frac{da^{(1)}}{dz^{(1)}} \frac{dz^{(2)}}{da^{(1)}} \frac{d\mathcal{L}}{dz^{(2)}}$$

We **can** calculate each of these Jacobians!

Example:

$$z^{(2)} = W^{(2)} a^{(1)}$$

$$\frac{dz^{(2)}}{da^{(1)}} = W^{(2)T}$$

Why might this be a **bad** idea?

if each $z^{(i)}$ or $a^{(i)}$ has about n dims...

each Jacobian is about $n \times n$ dimensions

matrix multiplication is $O(n^3)$

do we care?

AlexNet has layers with 4096 units...

Doing it more efficiently

this product is expensive
 this product is cheap: $O(n^2)$

$$\frac{d\mathcal{L}}{dW^{(1)}} = \frac{dz^{(1)}}{dW^{(1)}} \frac{da^{(1)}}{dz^{(1)}} \frac{dz^{(2)}}{da^{(1)}} \frac{d\mathcal{L}}{dz^{(2)}}$$

$n \times n$ $n \times 1$

this is **always** true because
the loss is scalar-valued!

Idea: start on the right

compute $\frac{dz^{(2)}}{da^{(1)}} \frac{d\mathcal{L}}{dz^{(2)}} = \delta$ first

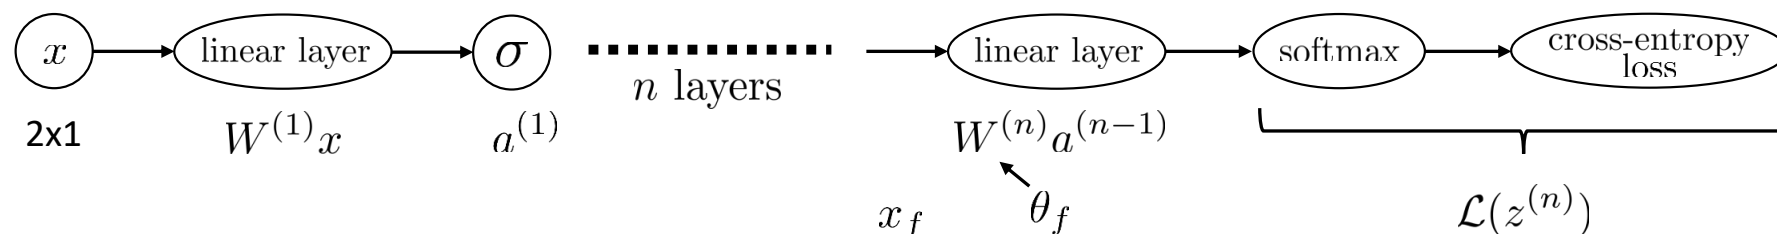

$$\frac{d\mathcal{L}}{dW^{(1)}} = \frac{dz^{(1)}}{dW^{(1)}} \underbrace{\frac{da^{(1)}}{dz^{(1)}} \delta}_{\text{this product is cheap: } O(n^2)}$$

compute $\frac{da^{(1)}}{dz^{(1)}} \delta = \gamma$

$$\frac{d\mathcal{L}}{dW^{(1)}} = \underbrace{\frac{dz^{(1)}}{dW^{(1)}} \gamma}_{\text{this product is cheap: } O(n^2)}$$

The backpropagation algorithm

“Classic” version



forward pass: calculate each $a^{(i)}$ and $z^{(i)}$ $a^{(n-1)} \xrightarrow{x_f, \theta_f} f \xrightarrow{} z^{(n-1)}$

backward pass:

initialize $\delta = \frac{d\mathcal{L}}{dz^{(n)}}$

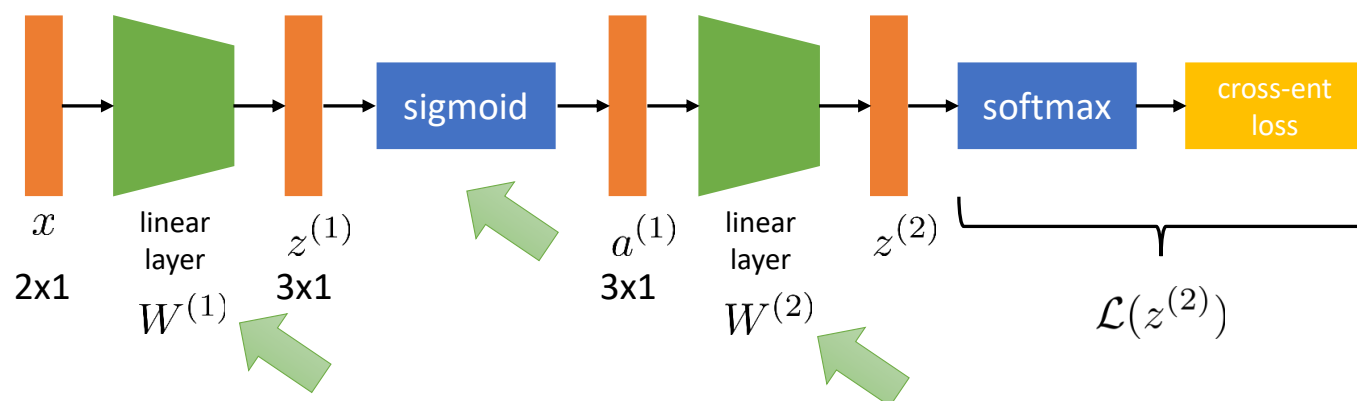
for each f with input x_f & params θ_f from end to start:

$$\frac{d\mathcal{L}}{d\theta_f} \leftarrow \frac{df}{d\theta_f} \delta$$

$$\delta \leftarrow \frac{df}{dx_f} \delta$$

Diagram illustrating the backpropagation of error gradients. A blue square represents the Jacobian of the function f with respect to its parameters θ_f . A green vertical bar represents the error gradient δ . The equation shows the blue square multiplied by the green bar equals another green bar, representing the backpropagation of the error gradient through the function f .

Let's walk through it...



forward pass: calculate each $a^{(i)}$ and $z^{(i)}$

backward pass:

→ initialize $\delta = \frac{d\mathcal{L}}{dz^{(n)}}$

for each f with input x_f & params θ_f from end to start:

→ $\frac{d\mathcal{L}}{d\theta_f} \leftarrow \frac{df}{d\theta_f} \delta$

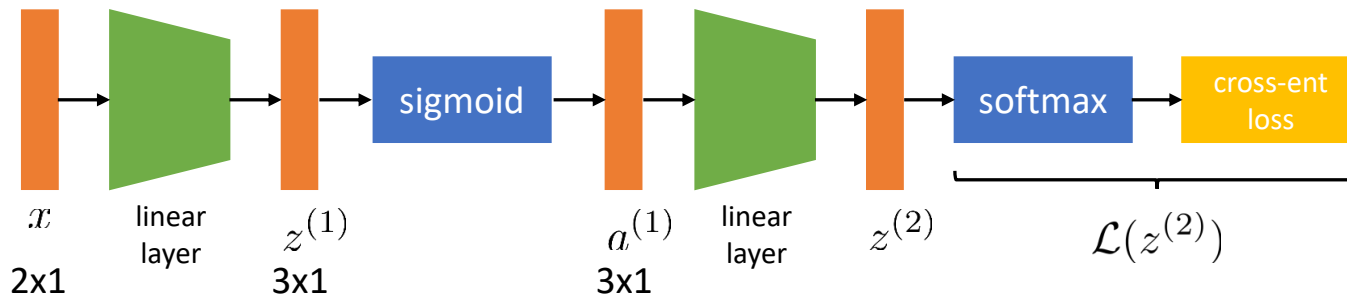
→ $\delta \leftarrow \frac{df}{dx_f} \delta$

$$\frac{d\mathcal{L}}{dW^{(2)}} = \underbrace{\frac{dz^{(2)}}{dW^{(2)}}}_{\delta} \frac{d\mathcal{L}}{dz^{(2)}}$$

$$\frac{d\mathcal{L}}{dW^{(1)}} = \underbrace{\frac{dz^{(1)}}{dW^{(1)}} \frac{da^{(1)}}{dz^{(1)}} \frac{dz^{(2)}}{da^{(1)}} \frac{d\mathcal{L}}{dz^{(2)}}}_{\delta}$$

Practical implementation

Neural network architecture details




Some things we should figure out:

How many layers?


How big are the layers?

What type of **activation function**?

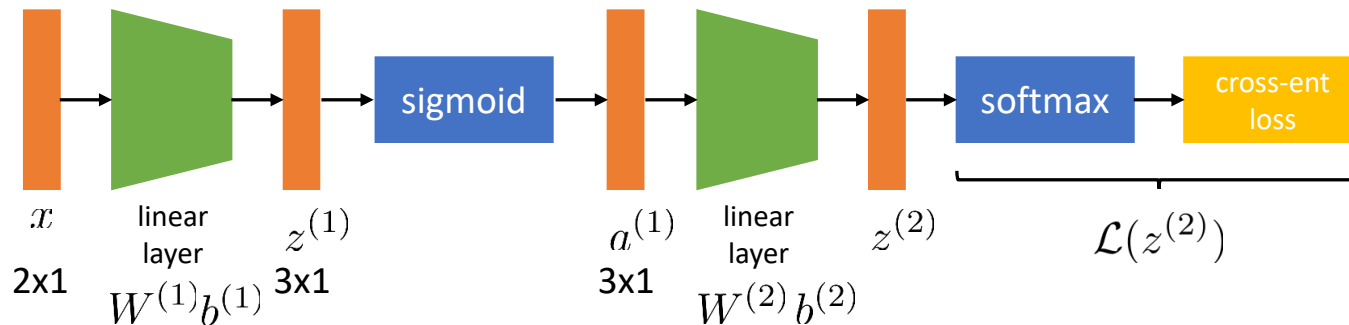
$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$



$$\text{ReLU}(x) = \max(0, x)$$



Bias terms



Linear layer:

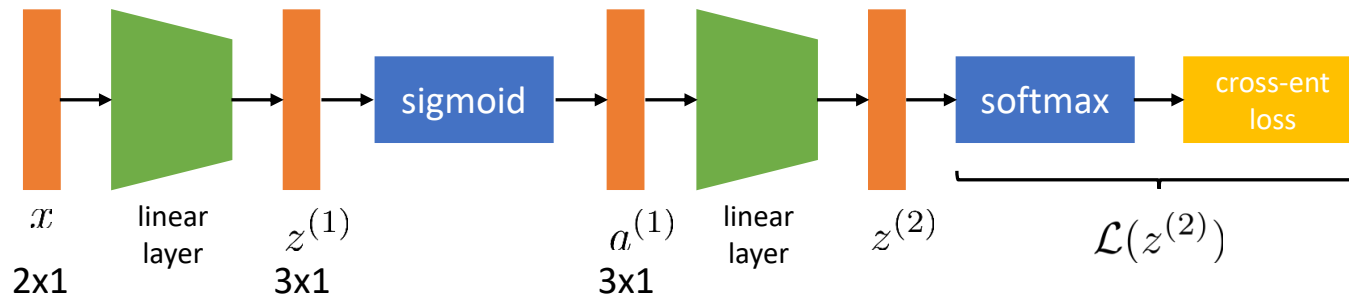
$$z^{(i+1)} = W^{(i)} a^{(i)} \quad \text{problem: if } a^{(i)} = \vec{0}, \text{ we always get } 0 \dots$$

Solution: add a "bias": has nothing to do with bias/variance bias

$$z^{(i+1)} = W^{(i)} a^{(i)} + b^{(i)}$$

additional parameters in each linear layer

What else do we need for backprop?



forward pass: calculate each $a^{(i)}$ and $z^{(i)}$

for each function, we need to compute:

backward pass:

initialize $\delta = \frac{d\mathcal{L}}{dz^{(n)}}$

for each f with input x_f & params θ_f from end to start:

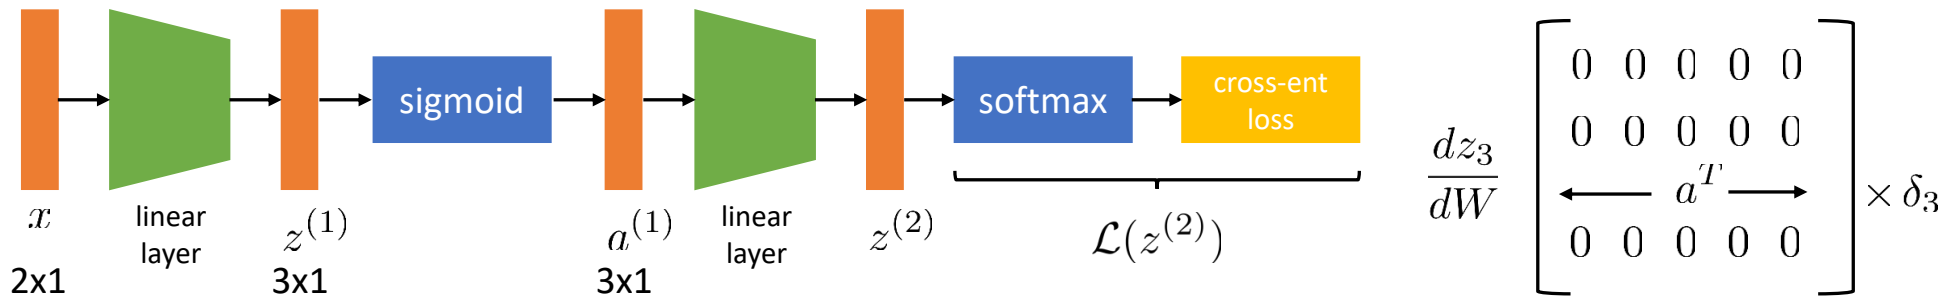
$$\frac{d\mathcal{L}}{d\theta_f} \leftarrow \frac{df}{d\theta_f} \delta$$

$$\delta \leftarrow \frac{df}{dx_f} \delta$$

$$\frac{df}{d\theta_f} \delta \quad \frac{df}{dx_f} \delta$$

linear layer
softmax + cross-entropy
sigmoid
ReLU

Backpropagation recipes: linear layer

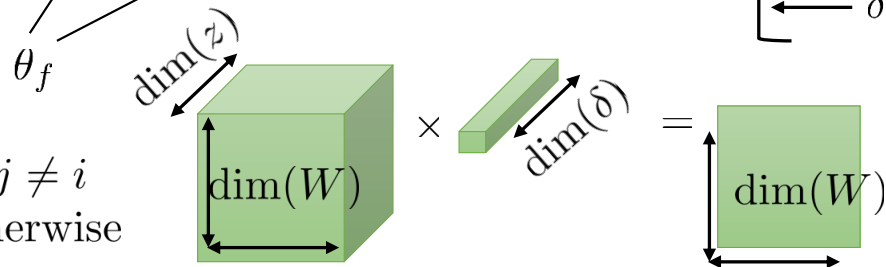


for each function, we need to compute: $\frac{df}{d\theta_f} \delta$ $\frac{df}{dx_f} \delta$ x_f

linear layer: $z^{(i+1)} = W^{(i)} a^{(i)} + b^{(i)}$ $z = Wa + b$ (just to simplify notation!)

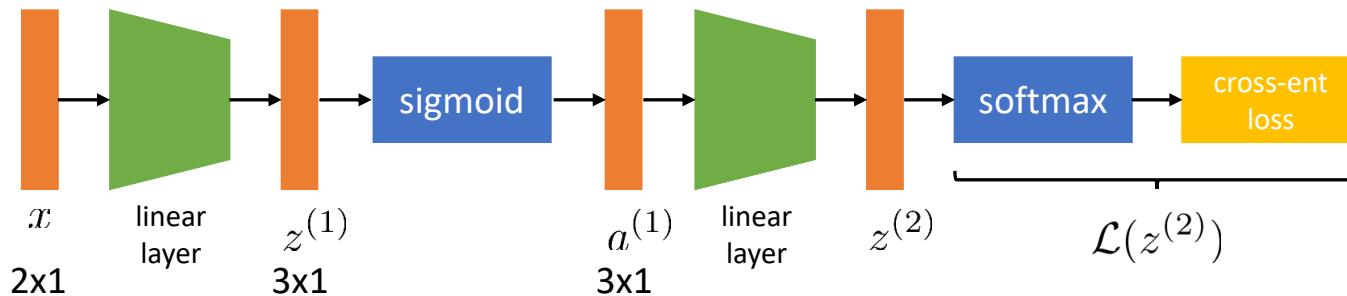
$$\frac{dz}{dW} \delta = \sum_i \frac{dz_i}{dW} \delta_i = \delta a^T$$

$$z_i = \sum_k W_{ik} a_k + b_i \quad \frac{dz_i}{dW_{jk}} = \begin{cases} 0 & \text{if } j \neq i \\ a_k & \text{otherwise} \end{cases}$$



$$\begin{bmatrix} \leftarrow \delta_1 a^T \rightarrow \\ \leftarrow \delta_2 a^T \rightarrow \\ \leftarrow \delta_3 a^T \rightarrow \\ \leftarrow \delta_4 a^T \rightarrow \end{bmatrix}$$

Backpropagation recipes: linear layer



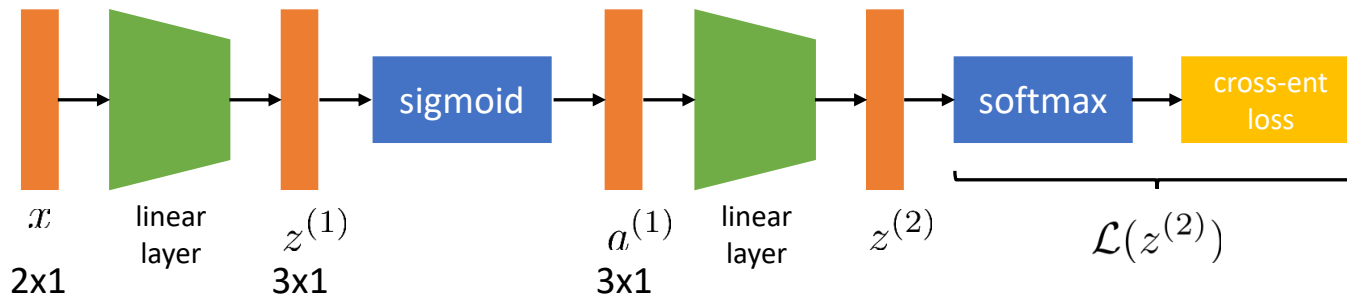
for each function, we need to compute: $\frac{df}{d\theta_f} \delta$ $\frac{df}{dx_f} \delta$

linear layer: $z^{(i+1)} = W^{(i)} a^{(i)} + b^{(i)}$ $z = Wa + b$ (just to simplify notation!)

$$\frac{dz}{db} \delta = \delta$$

$$z_i = \sum_k W_{ik} a_k + b_i \quad \frac{dz_i}{db_j} = \text{Ind}(i = j) \quad \frac{dz}{db} = \mathbf{I}$$

Backpropagation recipes: linear layer



for each function, we need to compute: $\frac{df}{d\theta_f} \delta$ $\frac{df}{dx_f} \delta$


linear layer: $z^{(i+1)} = W^{(i)} a^{(i)} + b^{(i)}$ $z = Wa + b$ (just to simplify notation!)

$$\frac{dz}{da} \delta = W^T \delta$$

$$z_i = \sum_k W_{ik} a_k + b_i \quad \frac{dz_i}{da_k} = W_{ik} \quad \frac{dz}{da} = W^T$$

$$\left(\frac{dy}{dx} \right)_{ij} = \frac{dy_j}{dx_i}$$

Backpropagation recipes: linear layer

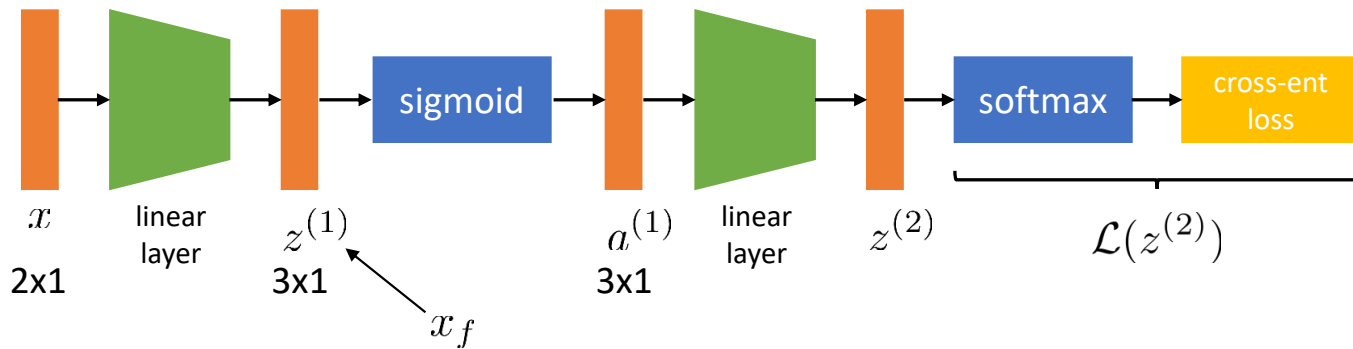


for each function, we need to compute: $\frac{df}{d\theta_f} \delta$ $\frac{df}{dx_f} \delta$

linear layer: $z^{(i+1)} = W^{(i)} a^{(i)} + b^{(i)}$ $z = Wa + b$ (just to simplify notation!)

$$\underbrace{\frac{dz}{da} \delta = W^T \delta}_{\frac{df}{dx_f} \delta} \quad \underbrace{\frac{dz}{dW} \delta = \delta a^T \quad \frac{dz}{db} \delta = \delta}_{\frac{df}{d\theta_f} \delta}$$

Backpropagation recipes: sigmoid



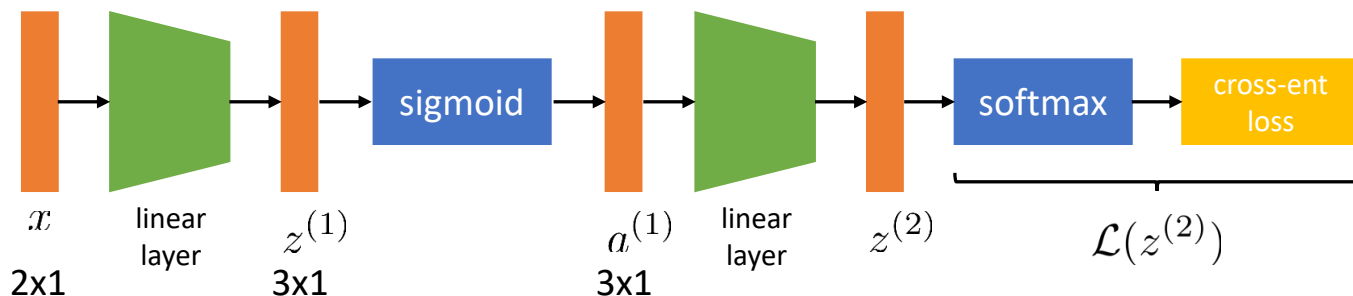
$$\frac{df}{dx} \begin{bmatrix} \frac{df_1}{dx_1} & 0 & 0 & 0 \\ 0 & \frac{df_2}{dx_2} & 0 & 0 \\ 0 & 0 & \frac{df_3}{dx_3} & 0 \\ 0 & 0 & 0 & \frac{df_4}{dx_4} \end{bmatrix}$$

for each function, we need to compute: $\frac{df}{d\theta_f} \delta$ $\frac{df}{dx_f} \delta$

$$\sigma(z_i) = \frac{1}{1 + \exp(-z_i)} \quad \frac{df_i}{dz_i} = \underbrace{\frac{\exp(-z_i)}{1 + \exp(-z_i)}}_{1 - \sigma(z_i)} \underbrace{\frac{1}{1 + \exp(-z_i)}}_{\sigma(z_i)} = (1 - \sigma(z_i))\sigma(z_i)$$

$$\left(\frac{df}{dz} \delta \right)_i = (1 - \sigma(z_i))\sigma(z_i)\delta_i \quad \underbrace{\frac{1 + \exp(-z_i)}{1 + \exp(-z_i)} - \frac{1}{1 + \exp(-z_i)}}_{1 - \sigma(z_i)} \sigma(z_i)$$

Backpropagation recipes: ReLU




for each function, we need to compute: $\frac{df}{d\theta_f} \delta$ $\frac{df}{dx_f} \delta$

$$f_i(z_i) = \max(0, z_i) \quad \frac{df_i}{dz_i} = \text{Ind}(z_i \geq 0)$$

$$\left(\frac{df}{dz} \delta \right)_i = \text{Ind}(z_i \geq 0) \delta_i$$

Summary



forward pass: calculate each $a^{(i)}$ and $z^{(i)}$

for each function, we need to compute:

backward pass:

initialize $\delta = \frac{d\mathcal{L}}{dz^{(n)}}$

for each f with input x_f & params θ_f from end to start:

$$\frac{d\mathcal{L}}{d\theta_f} \leftarrow \frac{df}{d\theta_f} \delta$$

$$\delta \leftarrow \frac{df}{dx_f} \delta$$

$$\frac{df}{d\theta_f} \delta \quad \frac{df}{dx_f} \delta$$

linear layer
softmax + cross-entropy
sigmoid
ReLU